

CS100433



Shadows

Junqiao Zhao 赵君峤

Department of Computer Science and Technology
College of Electronics and Information Engineering
Tongji University

Why shadows?


- 3D cueing
 - location, even motion



[Tom Thorne, Edinburgh]

How shadows are generated?

- Areas hidden from the light source by occlusion cause by objects
- Hard shadows and soft shadows




[Tom Thorne, Edinburgh]

Shadow techniques

- Ground shadow
- Shadow texture
- Shadow map
- Shadow volume

Ground shadow

- Shadow cast by objects onto the ground ($y=0$)




Point light ground shadows

- Blinn 88
- The matrix transform the object onto the ground

$$\begin{bmatrix} X_{sw} \\ 0 \\ Z_{sw} \\ 1 \end{bmatrix} = \begin{bmatrix} l_y & -l_x & 0 & 0 \\ 0 & 0 & 0 & 0 \\ 0 & -l_z & l_y & 0 \\ 0 & -1 & 0 & l_y \end{bmatrix} \begin{bmatrix} X_p \\ y_p \\ Z_p \\ 1 \end{bmatrix}$$


Draw the ground shadow

- The matrix transform the object onto the ground
- Thus
 - Draw the object
 - Multiply the shadow matrix
 - Redraw the object in grey





Problem of ground shadow

- The shadow only cast onto (ground) planes
- The shadow is hard shadow
- The performance is not optimal in static scene



Shadow texture

- Using a shadow image as a texture
 - Occluder from light's view
- Project the image onto the object
 - Can be curved surface or other objects




Shadow map



- Preparation
 - Prepare a depth buffer for each light
 - Render the scene from the light position
 - Save the depth information in the Depth buffer
- Rendering the scene
 - Render the objects; whenever rendering an object, check if it is shadowed or not by transforming its coordinate into the light space
 - After the transformation, if the depth value is larger than that in the light's depth buffer it should be shadowed

https://learnopengl.com/code_viewer_gh.php?code=src/5.advanced_lighting/3.1.2.shadow_mapping_base/shadow_mapping_base.cpp

Shadow map

- Using Depth buffer





[Wikipedia]




Figure 2: A conventional $2,048 \times 2,048$ pixel shadow map (left) compared to a 16 MB ASM (right).
EFFECTIVE SHADOW MAP SIZE: $65,536 \times 65,536$ PIXELS.


Shadow volume

- In the reality, the shadow cast by an object blocking light is a volume which is 3D!
- Project a ray from the light source through each vertex/silhouette in the shadow caster to infinity
- Any objects intersecting with the volume will get shadow on them.
 - Self-cast shadow
 - General-purpose


Shadow volumes

- Cast by silhouette





shadowed scene




wireframe shadow volumes


https://en.wikipedia.org/wiki/Shadow_volume

Stencil buffer

- Stencil buffer
 - A data buffer
 - Used as a Stencil
 - An integer per pixel
- Depth fail method
 - <http://ogldev.atspace.co.uk/www/tutorial40/tutorial40.html>
 - 1 Render the scene into Depth buffer
 - 2 Create shadow volume
 - 3 Render the volume in Stencil buffer following rules
 - 4 Render the scene under Stencil test



- Computation intensive
- Hard shadow




Stencil buffer based implementation

- Doom3 Made it popular



Soft shadows

- Made by area light
 - umbra – totally blocked from the light source
 - Penumbra – partially blocked from the light source
- Can be modelled by a collection of point light sources



Shadow map vs Shadow volume

- Faster than shadow volume
 - GPU based
- Less accurate because the resolution of the depth buffer
 - Aliasing at edges

- Questions?

CS100433

Ray tracing

Junqiao Zhao 赵君峤


Department of Computer Science and Technology
College of Electronics and Information Engineering
Tongji University

OpenGL vs Ray tracing

- OpenGL is based on a pipeline model in which primitives are rendered one at time
 - No shadows (shadow maps etc.)
 - No multiple reflections (environment maps etc.)
- Global approaches
 - Rendering equation
 - Ray tracing
 - Radiosity
 - Easily making effects: shadows, reflections, transparency


OpenGL vs Ray tracing

- Based on following light rays through the scene
- Iterate over pixels instead of primitives



Ray tracing



```
For each pixel {  
    Shoot a ray from camera to pixel;  
    //Perspective  
  
    for all objects in scene  
        Compute intersection with ray  
  
    Find object with closest intersection  
  
    Recursively shoot rays from the  
    intersection point  
  
    Display color using object + light  
    property  
}
```



(Ed Angel, NewMexico)


Ray tracing

- Sometimes the ray misses all of the objects
- and sometimes the ray will hit an object



Ray tracing


- If the ray hits an object, we want to know if that point on the object is in a shadow.
- So, when the ray hits an object, a secondary ray, called a "**shadow**" ray, is shot towards the light sources.



(Siggraph Education: Overview of Ray Tracing)

Ray tracing


- If this shadow ray hits another object before it hits a light source, then the first intersection point is in the shadow of the second object.
- For a simple illumination model this means that we only apply the ambient term for that light source.



(Siggraph Education: Overview of Ray Tracing)

Ray tracing


- If this shadow ray hits another object before it hits a light source, then the first intersection point is in the shadow of the second object.
- For a simple illumination model this means that we only apply the ambient term for that light source.



(Siggraph Education: Overview of Ray Tracing)

Ray tracing


- Also, when a ray hits an object, a **reflected ray** is generated which is tested against all of the objects in the scene.



(Siggraph Education: Overview of Ray Tracing)

Ray tracing


- If the reflected ray hits an object then a local illumination model is applied at the point of intersection and the result is carried back to the first intersection point.



(Siggraph Education: Overview of Ray Tracing)

Ray tracing

- If the intersected object is transparent, then a **transmitted ray** is generated and tested against all the objects in the scene.






(Siggraph Education: Overview of Ray Tracing)


Ray tracing

- The reflection ray can be implemented recursively
- There can be no reflection or multiple reflection

<http://www.siggraph.org/education/materials/HyperGraph/raytrace/rtrace0.htm>




Rays



(<http://www.ice.rwth-aachen.de/research/tools-projects/grace/ray-traycing/>)

Ray Tree

- Rays are generated recursively



Ray Tree





Fig. 12. The ray tree in schematic form.

Visualization of rays



<https://www.youtube.com/watch?v=aKqxonOrl4Q>

Ray tracing

- Should be able to handle all physical interactions
- Ray tracing paradigm is computation intensive
- Most rays do not affect what we see
- Scattering produces many (infinite) additional rays
- Alternative: ray casting


Diffuse Surfaces

- Theoretically the scattering at each point of intersection generates an infinite number of new rays that should be traced
- We could only trace the transmitted and reflected rays but use the Phong model to compute shade at point of intersection
- We can also sample the scattering rays by using Monte-carlo

Ray casting

- Only rays that reach the eye matter
- Reverse direction and cast rays
- Need at least one ray per pixel


```
For each pixel {  
    Shoot a ray from camera to pixel;  
    //Perspective  
  
    for all objects in scene  
        Compute intersection with ray  
  
    Find object with closest intersection  
  
    Display color using object + light  
    property  
}
```




(Ed Angel, NewMexico)

Ray casting for CSG

- CSG: Using set operations for solid shapes



http://escience.anu.edu.au/lecture/cg/surfaceModeling/CSG_tree.en.html




<http://www.cs.utah.edu/~shirley/fcg/figures/ray/>

Ray casting for CSG


- If only an image is needed do not need to change the models **Boolean operation**

- Ray casting
 - Find all intersections with a model, instead of just the closest.
 - find intervals where a ray is inside the object
 - Do set operations on intervals




Intersection and shading

- Ray intersection



Ray

- Standard representation: $r(t) = p + td$ where p is the start point and d is the direction and the parameter $t > 0$




Ray-sphere intersection

- Solving algebra:
- $r(t) = p + td$
- $(x - x_0)^2 + (y - y_0)^2 + (z - z_0)^2 = R^2$
- Substitute:
- $(r_x(t) - x_0)^2 + (r_y(t) - y_0)^2 + (r_z(t) - z_0)^2 = R^2$
- Quadratic formula to solve t
 - There will be two t values


Ray-box intersection

- Intersect with 6 faces
- Again solving algebra
 - Ray equation
 - Plane equation
 - Boundary check
- Or intersect with boundaries



Ray-box intersection

- Similar to clipping algorithm




Ray-box intersection

- Similar to clipping algorithm

- $p_x + t_{xmin}d_x = x_{min}$

- $p_x + t_{xmax}d_x = x_{max}$



Ray-box intersection


- Similar to clipping algorithm

- $p_x + t_{xmin}d_x = x_{min}$

- $p_x + t_{xmax}d_x = x_{max}$

- $p_y + t_{ymin}d_y = y_{min}$

- $p_y + t_{ymax}d_y = y_{max}$




Ray-box intersection

- Find the enter point and exit point

- $t_{xenter} = \min(t_{xmin}, t_{xmax})$
- $t_{yenter} = \min(t_{ymin}, t_{ymax})$
- $t_{xexit} = \max(t_{xmin}, t_{xmax})$
- $t_{yexit} = \max(t_{ymin}, t_{ymax})$


- Last enter point and first exit point

- $t_{enter} = \max(t_{xenter}, t_{yenter})$
- $t_{exit} = \min(t_{xexit}, t_{yexit})$



Shading

- Compute illumination
- Inputs:
 - Eye direction
 - Light source directions
 - Surface normal
 - Surface material
- Can be the same as Phong reflection model!
- $I = L_a k_a + L_d k_d \max(0, l \cdot n) + L_s k_s (r \cdot v)^\alpha$



Real-time ray tracing



<https://www.youtube.com/watch?v=aKqxonOrl4Q>

Ray tracing acceleration


- Ray tracer spends most of the time in ray-surface intersection
- To improve
 - Make intersection more efficient
 - Only do intersection when necessary
 - Trace in parallel
 - ...

Intersect when necessary

- Try to avoid meaningless intersections by using auxiliary data structures
 - BVH
 - Octree
 - BSPtree
 - Kdtree
 - Etc.


Bounding volume hierarchy (BVH)

- Object is contained within an volume
- First test the ray with the volume, then test the including objects when it hits the volume
- The volume should be simple (AABB boxes in most of time)
- Volume should fit the object tightly




Bounding volume hierarchy (BVH)

- Hierarchy of bounding volumes
- A tree




Bounding volume hierarchy (BVH)

- Hierarchy of bounding volumes
- A tree




Bounding volume hierarchy (BVH)

- Hierarchy of bounding volumes
- A tree



Bounding volume hierarchy (BVH)

- Hierarchy of bounding volumes
- A tree




Bounding volume hierarchy (BVH)

- Building a BVH
- Partition smartly
- Balance the tree will generally improve the efficiency
- Expectation of the intersection cost


Octree

- Analogy to Grids in 3D
- Not group objects, but divide space, how?



Octree

- Traverse the grid using enter point and exit point



Ray Tracing in One Weekend

- *Peter Shirley*

<https://raytracing.github.io/books/RayTracingInOneWeekend.html>

References

- Steve Marschner, CS4620/5620 Computer Graphics, Cornell
- Ed Angel, CS/EECE 433 Computer Graphics, University of New Mexico
- Elif Tosun, Computer Graphics, The University of New York

- Questions?